



# 什么是梯度

---

主讲人：龙良曲

# Clarification


- 导数, derivate
- 偏微分, partial derivate
- 梯度, gradient

$$\nabla f = \left( \frac{\partial f}{\partial x_1}; \frac{\partial f}{\partial x_2}; \dots; \frac{\partial f}{\partial x_n} \right)$$



$$z = y^2 - x^2$$

$$\frac{\partial z}{\partial x} = -2x$$

$$\frac{\partial z}{\partial y} = 2y$$



# What does grad mean?



# How to search for minima?

$$\theta_{t+1} = \theta_t - \alpha_t \nabla f(\theta_t).$$

Function:

$$J(\theta_1, \theta_2) = \theta_1^2 + \theta_2^2$$

Objective:

$$\min_{\theta_1, \theta_2} J(\theta_1, \theta_2)$$

Update rules:


$$\begin{aligned}\theta_1 &:= \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1, \theta_2) \\ \theta_2 &:= \theta_2 - \alpha \frac{d}{d\theta_2} J(\theta_1, \theta_2)\end{aligned}$$

Derivatives:

$$\frac{d}{d\theta_1} J(\theta_1, \theta_2) = \frac{d}{d\theta_1} \theta_1^2 + \frac{d}{d\theta_1} \theta_2^2 = 2\theta_1$$




$$\frac{d}{d\theta_2} J(\theta_1, \theta_2) = \frac{d}{d\theta_2} \theta_1^2 + \frac{d}{d\theta_2} \theta_2^2 = 2\theta_2$$

# Learning process-1







# Learning process-2



# Convex function





# Local Minima







# ResNet-56



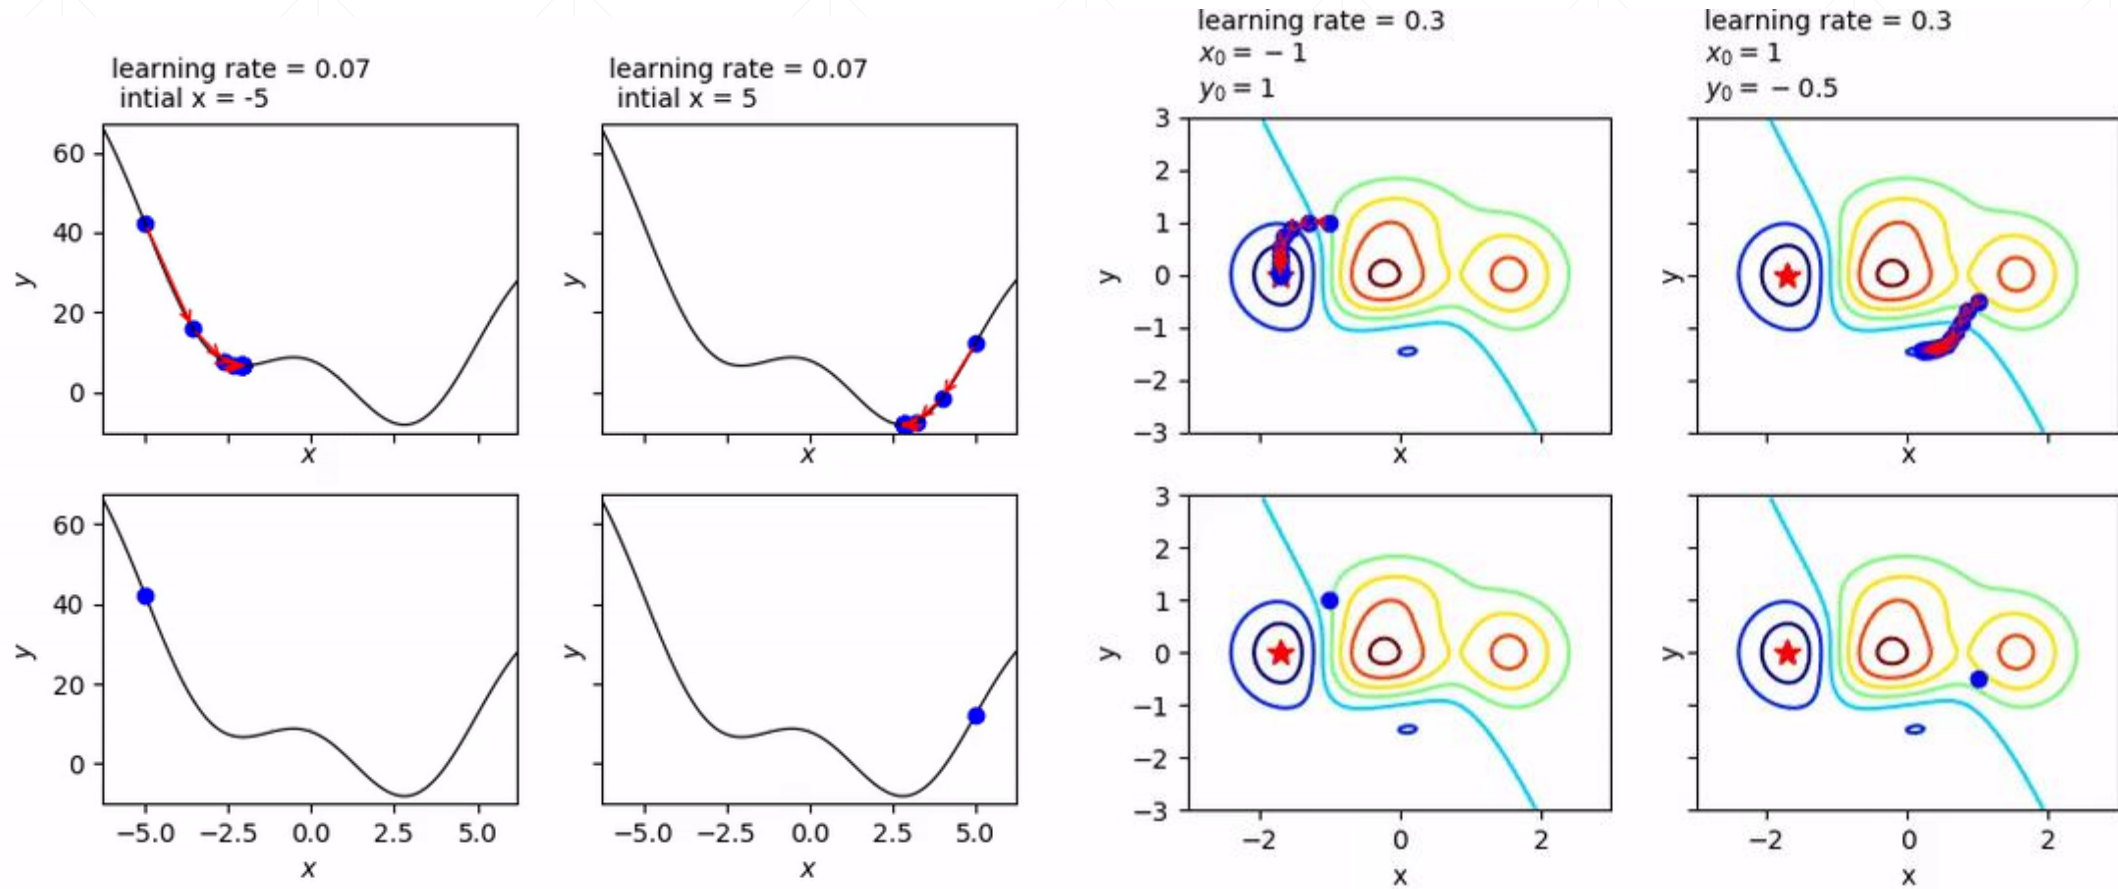
# Saddle point



# Optimizer Performance


- initialization status
  - learning rate
  - momentum
  - etc.
-

# Initialization







# Initialization



# Learning rate



# Escape minima



# 下一课时

---

常见函数梯度



**Thank You.**

---



# 常见函数梯度


---

主讲人：龙良曲

# Common Functions

Common Functions	Function	Derivative
Constant	$c$	$0$
Line	$x$	$1$
	$ax$	$a$
Square	$x^2$	$2x$
Square Root	$\sqrt{x}$	$(\frac{1}{2})x^{-\frac{1}{2}}$
Exponential	$e^x$	$e^x$
	$a^x$	$\ln(a) a^x$
Logarithms	$\ln(x)$	$1/x$
	$\log_a(x)$	$1 / (x \ln(a))$
Trigonometry (x is in <a href="#">radians</a> )	$\sin(x)$	$\cos(x)$
	$\cos(x)$	$-\sin(x)$
	$\tan(x)$	$\sec^2(x)$

$$xw + b$$






$$xw^2 + b^2$$




$$xe^w + e^b$$



$$[y - (xw + b)]^2$$



$$y \log(xw + b)$$





# 下一课时

---

什么是激活函数

**Thank You.**

---



# 激活函数及其梯度

---


主讲人：龙良曲

# Activation Functions





**PITTS WITH LETTVIN:** Pitts with Jerome Lettvin and one subject of their experiments on visual perception (1959).

Wikipedia




# Derivative



# Sigmoid / Logistic

$$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$$





# Derivative

$$\begin{aligned}\frac{d}{dx}\sigma(x) &= \frac{d}{dx} \left( \frac{1}{1 + e^{-x}} \right) \\&= \frac{e^{-x}}{(1 + e^{-x})^2} \\&= \frac{(1 + e^{-x}) - 1}{(1 + e^{-x})^2} \\&= \frac{1 + e^{-x}}{(1 + e^{-x})^2} - \left( \frac{1}{1 + e^{-x}} \right)^2 \\&= \sigma(x) - \sigma(x)^2 \\ \sigma' &= \sigma(1 - \sigma)\end{aligned}$$

---

# torch.sigmoid



```
In [5]: a=torch.linspace(-100,100,10)
```

```
In [6]: a
```

```
Out[6]:
```

```
tensor([-100.0000,  -77.7778,  -55.5556,  -33.3333,  -11.1111,   11.1111,  
         33.3333,   55.5555,   77.7778,  100.0000])
```

```
In [7]: torch.sigmoid(a)
```


```
Out[7]:
```

```
tensor([0.0000e+00,  1.6655e-34,  7.4564e-25,  3.3382e-15,  1.4945e-05,  9.9999e-01,  
        1.0000e+00,  1.0000e+00,  1.0000e+00,  1.0000e+00])
```

# Tanh

$$f(x) = \tanh(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$$

$$= 2\textcolor{red}{sigmoid}(2x) - 1$$



# Derivative

$$\begin{aligned}\frac{d}{dx} \tanh(x) &= \frac{(e^x + e^{-x})(e^x + e^{-x}) - (e^x - e^{-x})(e^x - e^{-x})}{(e^x + e^{-x})^2} \\ &= 1 - \frac{(e^x - e^{-x})^2}{(e^x + e^{-x})^2} = 1 - \tanh^2(x)\end{aligned}$$

---

# torch.tanh



```
In [9]: a=torch.linspace(-1,1,10)
```

```
In [10]: torch.tanh(a)
```

```
Out[10]:
```

```
tensor([-0.7616, -0.6514, -0.5047, -0.3215, -0.1107,  0.1107,  0.3215,  0.5047,  
        0.6514,  0.7616])
```


# Rectified Linear Unit

$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$



# Derivative

$$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$





# F.relu



```
In [11]: from torch.nn import functional as F
```

```
In [12]: a=torch.linspace(-1,1,10)
```

```
In [13]: torch.relu(a)
```

```
Out[13]:
```

```
tensor([0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.1111, 0.3333, 0.5556, 0.7778,  
        1.0000])
```

```
In [14]: F.relu(a)
```

```
Out[14]:
```

```
tensor([0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.1111, 0.3333, 0.5556, 0.7778,  
        1.0000])
```

# 下一课时

---

Loss及其梯度

**Thank You.**

---



# LOSS及其梯度

---

主讲人：龙良曲

# Typical Loss

- Mean Squared Error
  - Cross Entropy Loss
    - binary
    - multi-class
    - +softmax
    - Leave it to Logistic Regression Part
-


# MSE

- $\text{loss} = \sum [y - (xw + b)]^2$
  - $L2 - \text{norm} = \|y - (xw + b)\|_2$
  - $\text{loss} = \text{norm}(y - (xw + b))^2$
-

# Derivative

- $\text{loss} = \sum [y - f_{\theta}(x)]^2$

- $\frac{\nabla \text{loss}}{\nabla \theta} = 2 \sum [y - f_{\theta}(x)] * \frac{\nabla f_{\theta}(x)}{\nabla \theta}$





# autograd.grad



```
In [15]: x=torch.ones(1)
In [17]: w=torch.full([1],2)
In [19]: mse=F.mse_loss(torch.ones(1), x*w)
Out[20]: tensor(1.)

In [21]: torch.autograd.grad(mse,[w])
#RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn

In [22]: w.requires_grad_()
Out[22]: tensor([2.], requires_grad=True)

In [23]: torch.autograd.grad(mse,[w])
#RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn

In [24]: mse=F.mse_loss(torch.ones(1), x*w)

In [25]: torch.autograd.grad(mse,[w])
Out[25]: (tensor([2.]),)
```

# loss.backward



```
In [15]: x=torch.ones(1)
In [17]: w=torch.full([1],2)
In [19]: mse=F.mse_loss(torch.ones(1), x*w)
Out[20]: tensor(1.)
```

```
In [21]: torch.autograd.grad(mse,[w])
#RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn
```

```
In [22]: w.requires_grad_()
Out[22]: tensor([2.], requires_grad=True)
```

```
In [23]: torch.autograd.grad(mse,[w])
#RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn
```

```
In [24]: mse=F.mse_loss(torch.ones(1), x*w)
In [27]: mse.backward()
```


```
In [28]: w.grad
Out[28]: tensor([2.])
```

# Gradient API

- `torch.autograd.grad(loss, [w1, w2,...])`
    - `[w1.grad, w2.grad,...]`
  - `loss.backward()`
    - `w1.grad`
    - `w2.grad`
-

# Softmax

- soft version of max



# Derivative

$$p_i = \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}$$

when  $i = j$

$$\frac{\partial p_i}{\partial a_j} = \frac{\partial \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}}{\partial a_j}$$

$$f(x) = \frac{g(x)}{h(x)}$$

$$f'(x) = \frac{g'(x)h(x) - h'(x)g(x)}{h(x)^2}$$

$$g(x) = e^{a_i}$$

$$h(x) = \sum_{k=1}^N e^{a_k}$$

$$\begin{aligned} \frac{\partial \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}}{\partial a_j} &= \frac{e^{a_i} \sum_{k=1}^N e^{a_k} - e^{a_j} e^{a_i}}{\left(\sum_{k=1}^N e^{a_k}\right)^2} \\ &= \frac{e^{a_i} \left(\sum_{k=1}^N e^{a_k} - e^{a_j}\right)}{\left(\sum_{k=1}^N e^{a_k}\right)^2} \\ &= \frac{e^{a_j}}{\sum_{k=1}^N e^{a_k}} \times \frac{\left(\sum_{k=1}^N e^{a_k} - e^{a_j}\right)}{\sum_{k=1}^N e^{a_k}} \\ &= p_i(1 - p_j) \end{aligned}$$

# Derivative

$$p_i = \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}$$

$$\frac{\partial p_i}{\partial a_j} = \frac{\partial \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}}{\partial a_j}$$

$$f(x) = \frac{g(x)}{h(x)}$$

$$f'(x) = \frac{g'(x)h(x) - h'(x)g(x)}{h(x)^2}$$

$$g(x) = e^{a_i}$$

$$h(x) = \sum_{k=1}^N e^{a_k}$$

when  $i \neq j$

$$\frac{\partial \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}}{\partial a_j} = \frac{0 - e^{a_j} e^{a_i}}{\left(\sum_{k=1}^N e^{a_k}\right)^2}$$

$$\begin{aligned} &= \frac{-e^{a_j}}{\sum_{k=1}^N e^{a_k}} \times \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}} \\ &= -p_j \cdot p_i \end{aligned}$$

# Derivative

$$\frac{\partial p_i}{\partial a_j} = \begin{cases} p_i(1 - p_j) & \text{if } i = j \\ -p_j \cdot p_i & \text{if } i \neq j \end{cases}$$

Or using Kronecker delta  $\delta_{ij} = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$

$$\frac{\partial p_i}{\partial a_j} = p_i(\delta_{ij} - p_j)$$

---



# F.softmax

```
In [29]: a=torch.rand(3) # tensor([0.1440, 0.5349, 0.7022])
```

```
In [33]: a.requires_grad_()
```

```
Out[33]: tensor([0.1440, 0.5349, 0.7022], requires_grad=True)
```

```
In [34]: p=F.softmax(a,dim=0)
```

```
In [35]: p.backward()
```

RuntimeError: Trying to backward through the graph a second time, but the buffers have already been freed. Specify retain\_graph=True when calling backward the first time.

```
In [38]: p=F.softmax(a,dim=0)
```

```
In [39]: torch.autograd.grad(p[1],[a],retain_graph=True)
```

```
Out[39]: (tensor([-0.0828, 0.2274, -0.1447]),)
```

```
In [40]: torch.autograd.grad(p[2],[a])
```

```
Out[40]: (tensor([-0.0979, -0.1447, 0.2425]),)
```



# 下一课时

---

## 链式法则

**Thank You.**

---