



Logistic Regression

主讲人：龙良曲

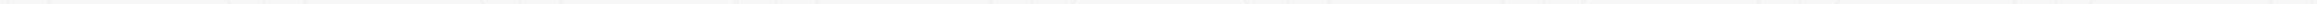
Recap

- for continuous: $y = xw + b$
- for probability output: $y = \sigma(xw + b)$
 - σ : *sigmoid or logistic*



Binary Classification

- interpret network as $f: x \rightarrow p(y|x; \theta)$
- output $\in [0, 1]$
- which is exactly what *logistic function* comes in!



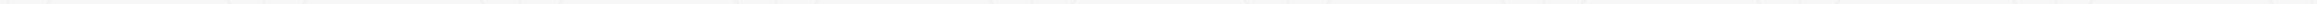
Goal v.s. Approach

- For regression:
 - Goal: $\text{pred} = y$
 - Approach: minimize $\text{dist}(\text{pred}, y)$
- For classification:
 - Goal: maximize benchmark, e.g. *accuracy*
 - Approach1: minimize $\text{dist}(p_\theta(y|x), p_r(y|x))$
 - Approach2: minimize $\text{divergence}(p_\theta(y|x), p_r(y|x))$



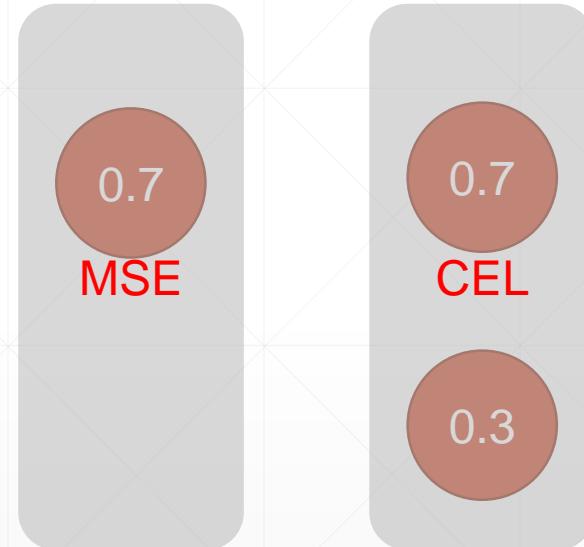
Q1. why not maximize accuracy?

- $acc. = \frac{\sum I(pred_i == y_i)}{len(Y)}$
- issues 1. gradient = 0 if accuracy unchanged but weights changed
- issues 2. *gradient not continuous* since the number of correct is not continuous



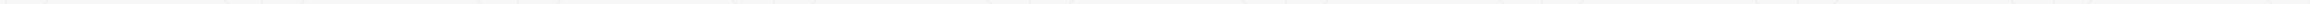
Q2. why call logistic regression

- use sigmoid
- Controversial!
 - MSE => regression
 - Cross Entropy => classification



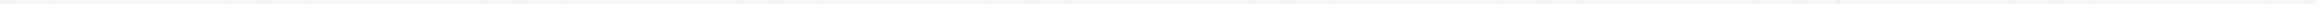
Binary Classification

- $f: x \rightarrow p(y = 1|x)$
 - if $p(y = 1|x) > 0.5$, predict as **1**
 - else predict as **0**
- minimize MSE
- confused?
 - <http://www.fharrell.com/post/classification/>



Multi-class classification

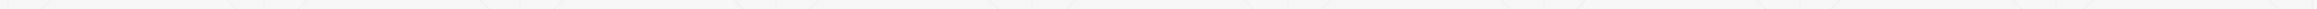
- $f: x \rightarrow p(y|x)$
 - $[p(y = 0|x), p(y = 1|x), \dots, p(y = 9|x)]$
- $p(y|x) \in [0, 1]$
- $\sum_{i=0}^9 p(y = i|x) = 1$



Softmax

- $\sum_{i=0}^9 p(y = i|x) = 1$

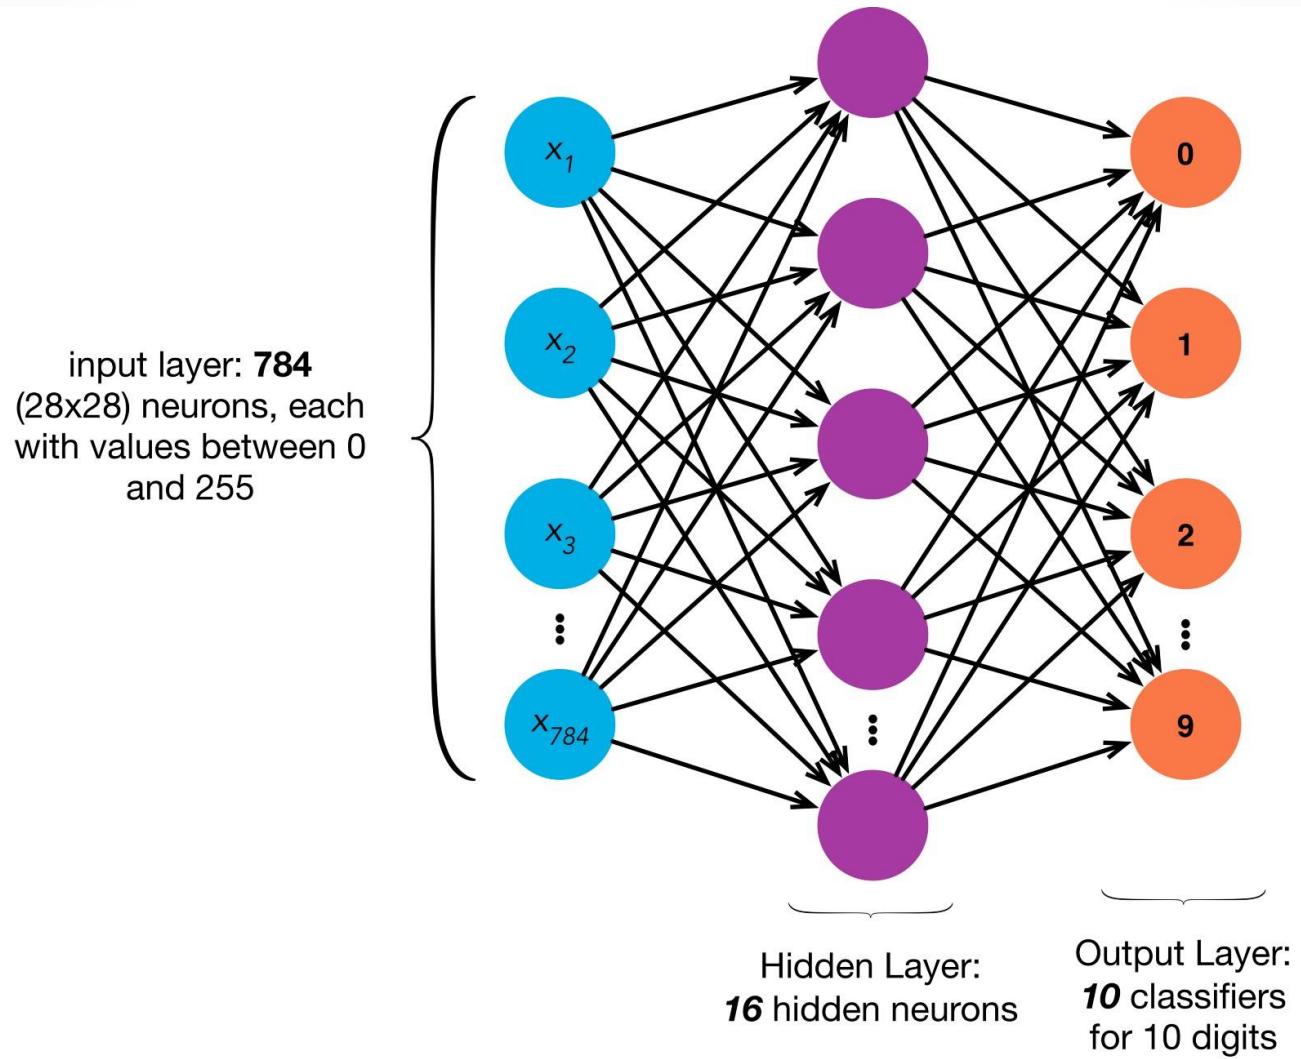
$$p_i = \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}$$



Softmax

- enlarger the larger

$$p_i = \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}$$



下一课时

交叉熵

Thank You.



交叉熵

主讲人：龙良曲

Why not MSE?

Label	predict	correct
3	[0.3, 0.3, 0.4]	yes
2	[0.3, 0.4, 0.3]	yes
1	[0.1, 0.2, 0.7]	no

0.54

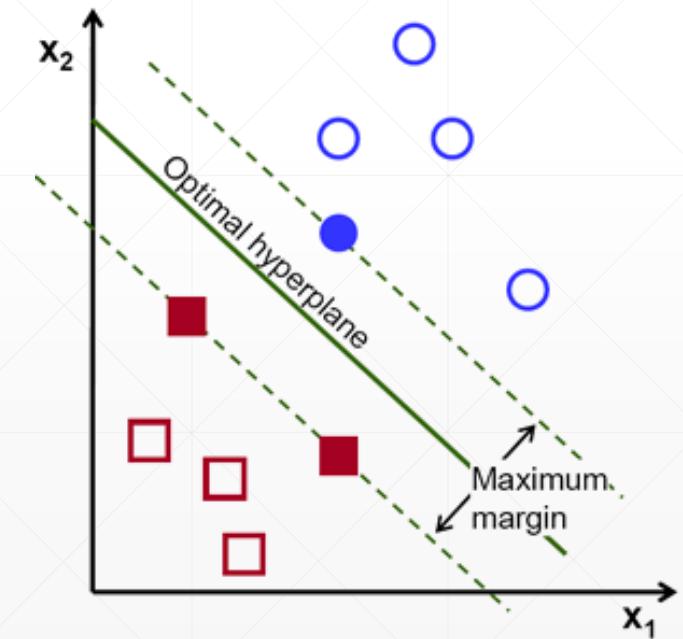
Label	predict	correct
3	[0.1, 0.2, 0.7]	yes
2	[0.1, 0.7, 0.2]	yes
1	[0.3, 0.4, 0.3]	no

0.34

Loss for classification

- MSE
- Cross Entropy Loss
- Hinge Loss

$$\sum_i \max(0, 1 - y_i * h_\theta(x_i))$$



Entropy

- Uncertainty
- measure of surprise
- higher entropy: **higher uncertainty.**

$$\text{Entropy} = - \sum_i P(i) \log P(i)$$



Claude Shannon

Lottery



```
In [2]: a=torch.full([4],1/4.)  
tensor([0.2500, 0.2500, 0.2500, 0.2500])
```

```
In [4]: a*torch.log2(a)  
Out[4]: tensor([-0.5000, -0.5000, -0.5000, -0.5000])
```

```
In [6]: -(a*torch.log2(a)).sum()  
Out[6]: tensor(2.)
```

```
In [7]: a=torch.tensor([0.1,0.1,0.1,0.7])  
In [8]: -(a*torch.log2(a)).sum()  
Out[8]: tensor(1.3568)
```

```
In [15]: a=torch.tensor([0.001,0.001,0.001,0.999])  
In [16]: -(a*torch.log2(a)).sum()  
Out[16]: tensor(0.0313)
```



Cross Entropy

$$H(p, q) = - \sum p(x) \log q(x)$$

$$H(p, q) = H(p) + D_{\text{KL}}(p|q).$$

- P=Q
 - cross Entropy = Entropy
- for one-hot encoding,
 - entropy = $1 \log 1 = 0$



Binary Classification

$$H(P, Q) = -P(cat) \log Q(cat) - (1 - P(cat)) \log(1 - Q(cat))$$

$$P(dog) = (1 - P(cat))$$

$$\begin{aligned} H(P, Q) &= - \sum_{i=(cat,dog)} P(i) \log Q(i) \\ &= -P(cat) \log Q(cat) - P(dog) \log Q(dog) \end{aligned}$$

$$-(y \log(p) + (1 - y) \log(1 - p))$$

for example



$$P_1 = [1 \ 0 \ 0 \ 0 \ 0]$$

$$Q_1 = [0.4 \ 0.3 \ 0.05 \ 0.05 \ 0.2]$$

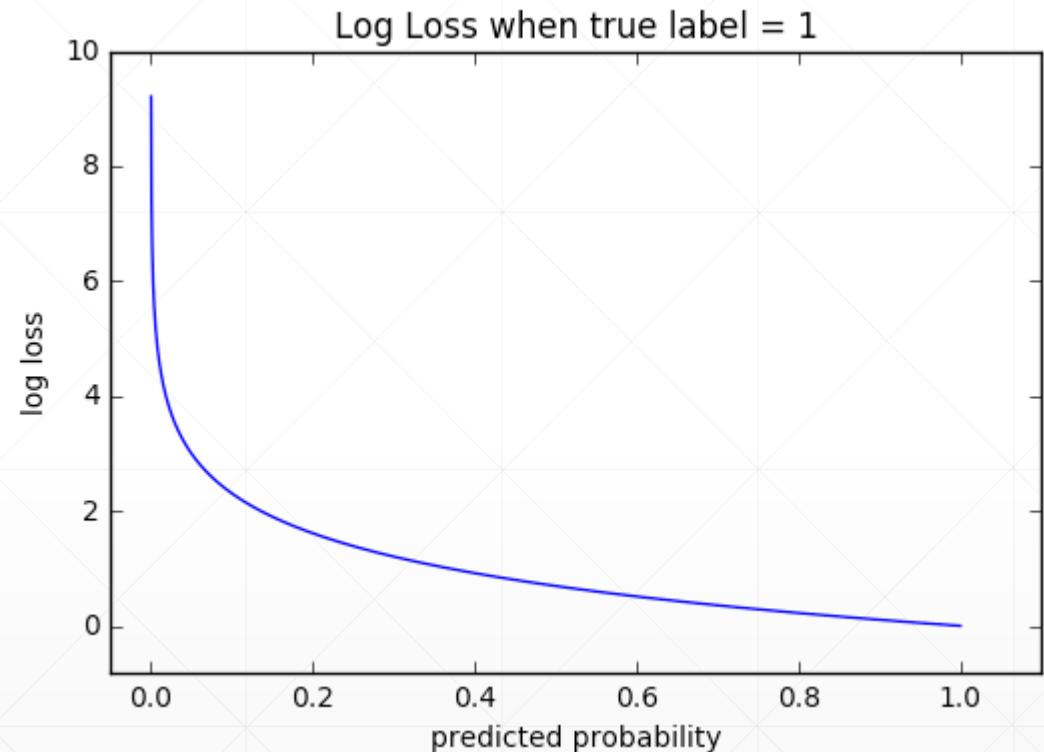
$$\begin{aligned} H(P_1, Q_1) &= - \sum_i P_1(i) \log Q_1(i) \\ &= -(1 \log 0.4 + 0 \log 0.3 + 0 \log 0.05 + 0 \log 0.05 + 0 \log 0.2) \\ &= -\log 0.4 \\ &\approx 0.916 \end{aligned}$$

$$Q_1 = [0.98 \ 0.01 \ 0 \ 0 \ 0.01]$$

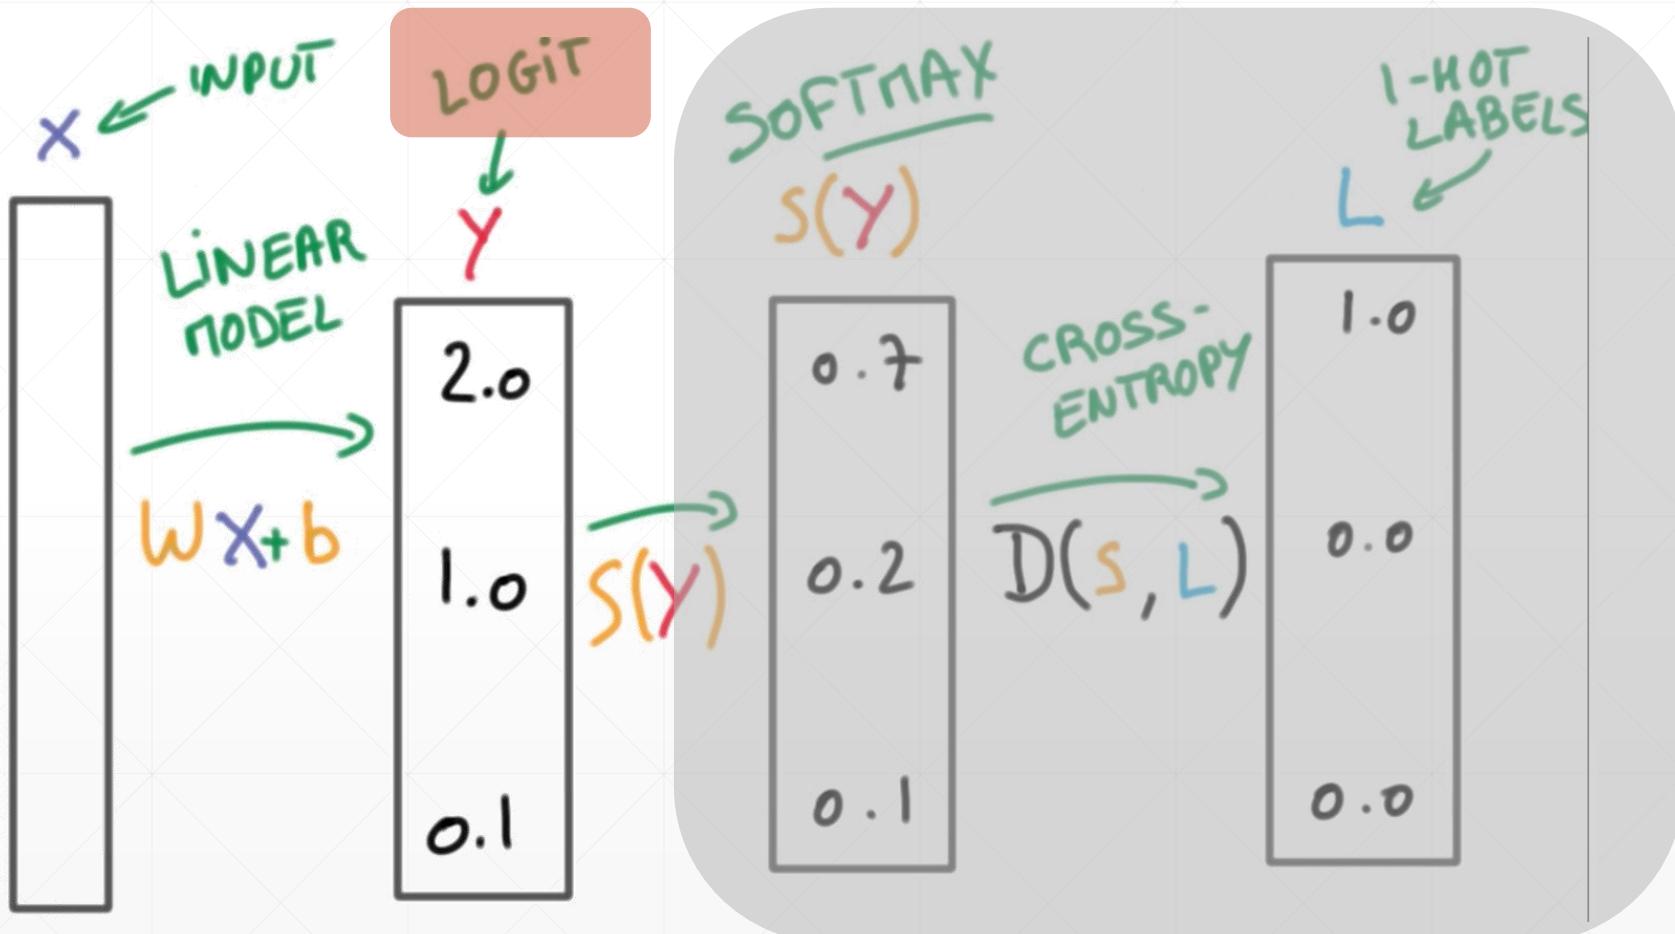
$$\begin{aligned} H(P_1, Q_1) &= - \sum_i P_1(i) \log Q_1(i) \\ &= -(1 \log 0.98 + 0 \log 0.01 + 0 \log 0 + 0 \log 0 + 0 \log 0.01) \\ &= -\log 0.98 \\ &\approx 0.02 \end{aligned}$$

why not use MSE

- sigmoid + MSE
 - gradient vanish
- converge slower
- But, sometimes
 - e.g. meta-learning



Therefore



Numerical Stability



```
In [17]: x=torch.randn(1,784)
```

```
In [18]: w=torch.randn(10,784)
```

```
In [19]: logits=x@w.t()
```

```
Out[20]: torch.Size([1, 10])
```

```
In [34]: pred=F.softmax(logits, dim=1)
```

```
Out[37]: torch.Size([1, 10])
```

```
In [40]: pred_log=torch.log(pred)
```

```
In [41]: F.cross_entropy(logits,torch.tensor([3]))
```

```
Out[41]: tensor(82.4905)
```

```
In [42]: F.nll_loss(pred_log,torch.tensor([3]))
```

```
Out[42]: tensor(82.4905)
```

下一课时

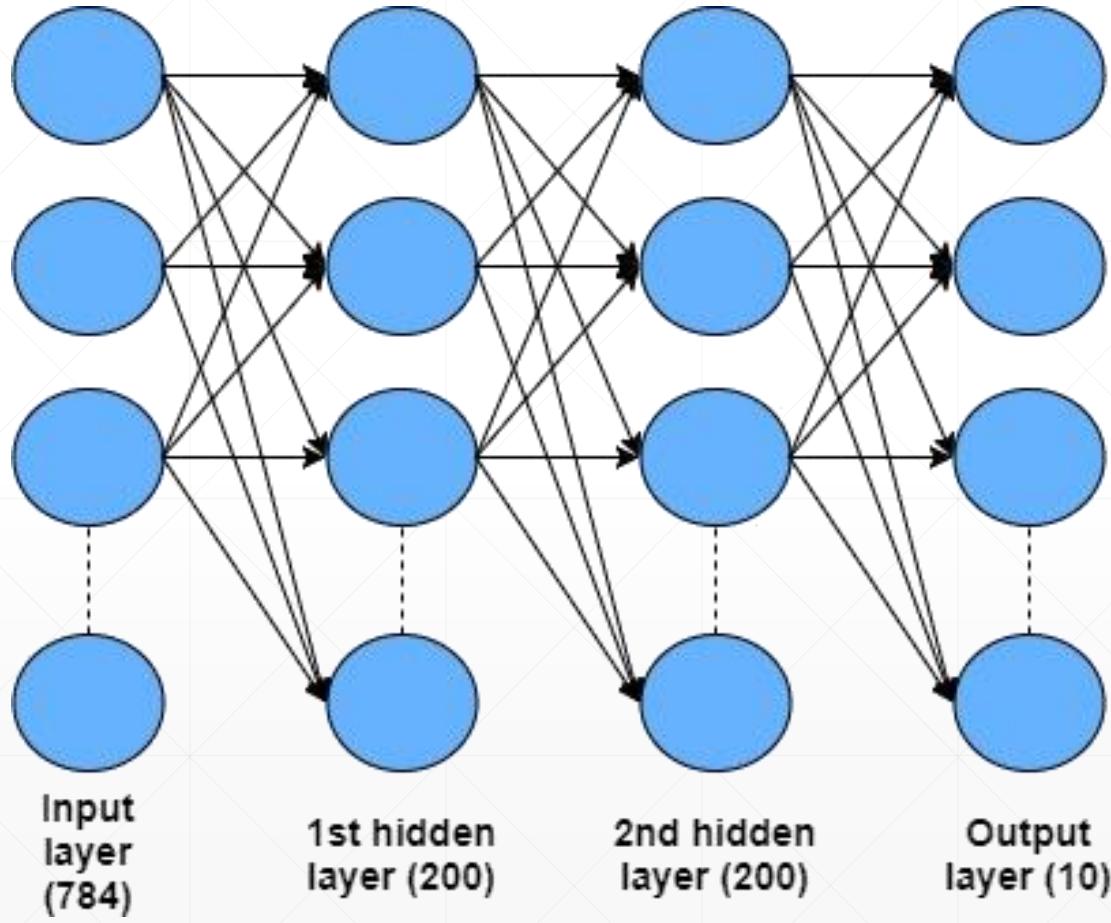
实战多分类问题

Thank You.



多分类问题

主讲人：龙良曲



Network Architecture

```
● ● ●

w1, b1 = torch.randn(200, 784, requires_grad=True), \
          torch.zeros(200, requires_grad=True)
w2, b2 = torch.randn(200, 200, requires_grad=True), \
          torch.zeros(200, requires_grad=True)
w3, b3 = torch.randn(10, 200, requires_grad=True), \
          torch.zeros(10, requires_grad=True)

def forward(x):
    x = x@w1.t() + b1
    x = F.relu(x)
    x = x@w2.t() + b2
    x = F.relu(x)
    x = x@w3.t() + b3
    x = F.relu(x)
    return x
```

Train

```
● ● ●  
  
optimizer = optim.SGD([w1, b1, w2, b2, w3, b3], lr=learning_rate)  
criterion = nn.CrossEntropyLoss()  
  
for epoch in range(epochs):  
  
    for batch_idx, (data, target) in enumerate(train_loader):  
        data = data.view(-1, 28*28)  
  
        logits = forward(data)  
        loss = criterion(logits, target)  
  
        optimizer.zero_grad()  
        loss.backward()  
        # print(w1.grad.norm(), w2.grad.norm())  
        optimizer.step()
```

```
C:\ProgramData\conda\python.exe F:/PytorchTutorial/lesson26-LR多分类实战/main.py
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Processing...
Done!
```

em....

```
● ● ●

w1, b1 = torch.randn(200, 784, requires_grad=True), \
          torch.zeros(200, requires_grad=True)
w2, b2 = torch.randn(200, 200, requires_grad=True), \
          torch.zeros(200, requires_grad=True)
w3, b3 = torch.randn(10, 200, requires_grad=True), \
          torch.zeros(10, requires_grad=True)

torch.nn.init.kaiming_normal_(w1)
torch.nn.init.kaiming_normal_(w2)
torch.nn.init.kaiming_normal_(w3)
```

下一课时

PyTorch全连接 层

Thank You.



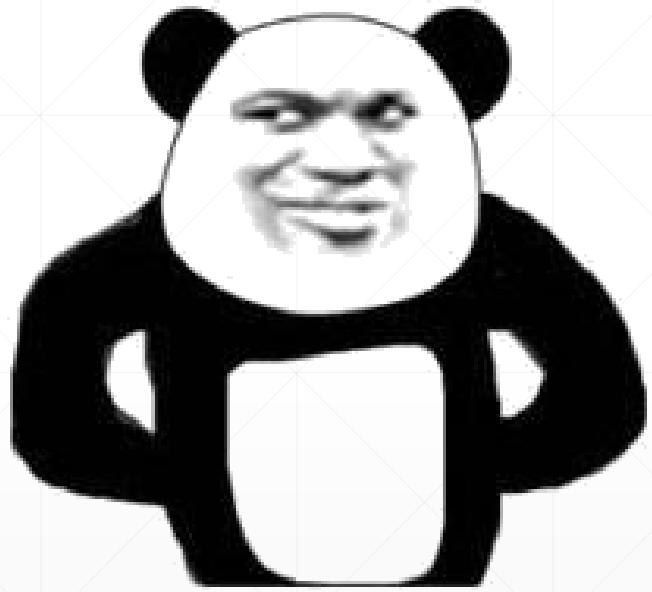
全军出击：全连接层

主讲人：龙良曲

I know nothing



Be practical



在大佬看来说这么多没用
还是直接上去猛干就对了



nn.Linear



```
In [44]: x.shape  
Out[44]: torch.Size([1, 784])
```

```
In [46]: layer1=nn.Linear(784, 200)  
In [47]: layer2=nn.Linear(200,200)  
In [48]: layer3=nn.Linear(200,10)
```

```
In [49]: x=layer1(x)  
In [50]: x.shape  
Out[50]: torch.Size([1, 200])
```

```
In [52]: x=layer2(x)  
In [53]: x.shape  
Out[53]: torch.Size([1, 200])
```

```
In [54]: x=layer3(x)  
In [55]: x.shape  
Out[55]: torch.Size([1, 10])
```

relu?



```
In [49]: x=layer1(x)
In [56]: x=F.relu(x, inplace=True)
In [50]: x.shape
Out[50]: torch.Size([1, 200])

In [52]: x=layer2(x)
In [56]: x=F.relu(x, inplace=True)
In [53]: x.shape
Out[53]: torch.Size([1, 200])

In [54]: x=layer3(x)
In [56]: x=F.relu(x, inplace=True)
In [55]: x.shape
Out[55]: torch.Size([1, 10])
```

concisely

- inherit from nn.Module
 - init layer in `__init__`
 - implement `forward()`
-

Step1.



```
class MLP(nn.Module):  
  
    def __init__(self):  
        super(MLP, self).__init__()
```



Step2.

```
class MLP(nn.Module):

    def __init__(self):
        super(MLP, self).__init__()

        self.model = nn.Sequential(
            nn.Linear(784, 200),
            nn.ReLU(inplace=True),
            nn.Linear(200, 200),
            nn.ReLU(inplace=True),
            nn.Linear(200, 10),
            nn.ReLU(inplace=True),
        )
```

Step3.

```
class MLP(nn.Module):

    def __init__(self):
        super(MLP, self).__init__()

        self.model = nn.Sequential(
            nn.Linear(784, 200),
            nn.ReLU(inplace=True),
            nn.Linear(200, 200),
            nn.ReLU(inplace=True),
            nn.Linear(200, 10),
            nn.ReLU(inplace=True),
        )

    def forward(self, x):
        x = self.model(x)

        return x
```

nn.ReLU v.s. F.relu()

- class-style API
- function-style API



```
In [55]: x.shape  
Out[55]: torch.Size([1, 10])  
  
In [56]: x=F.relu(x, inplace=True)  
  
In [57]: layer=nn.ReLU()  
  
In [58]: x=layer(x)
```

Train

```
● ● ●

net = MLP()
optimizer = optim.SGD(net.parameters(), lr=learning_rate)
criteon = nn.CrossEntropyLoss()

for epoch in range(epochs):

    for batch_idx, (data, target) in enumerate(train_loader):
        data = data.view(-1, 28*28)

        logits = net(data)
        loss = criteon(logits, target)

        optimizer.zero_grad()
        loss.backward()
        # print(w1.grad.norm(), w2.grad.norm())
        optimizer.step()
```

下一课时

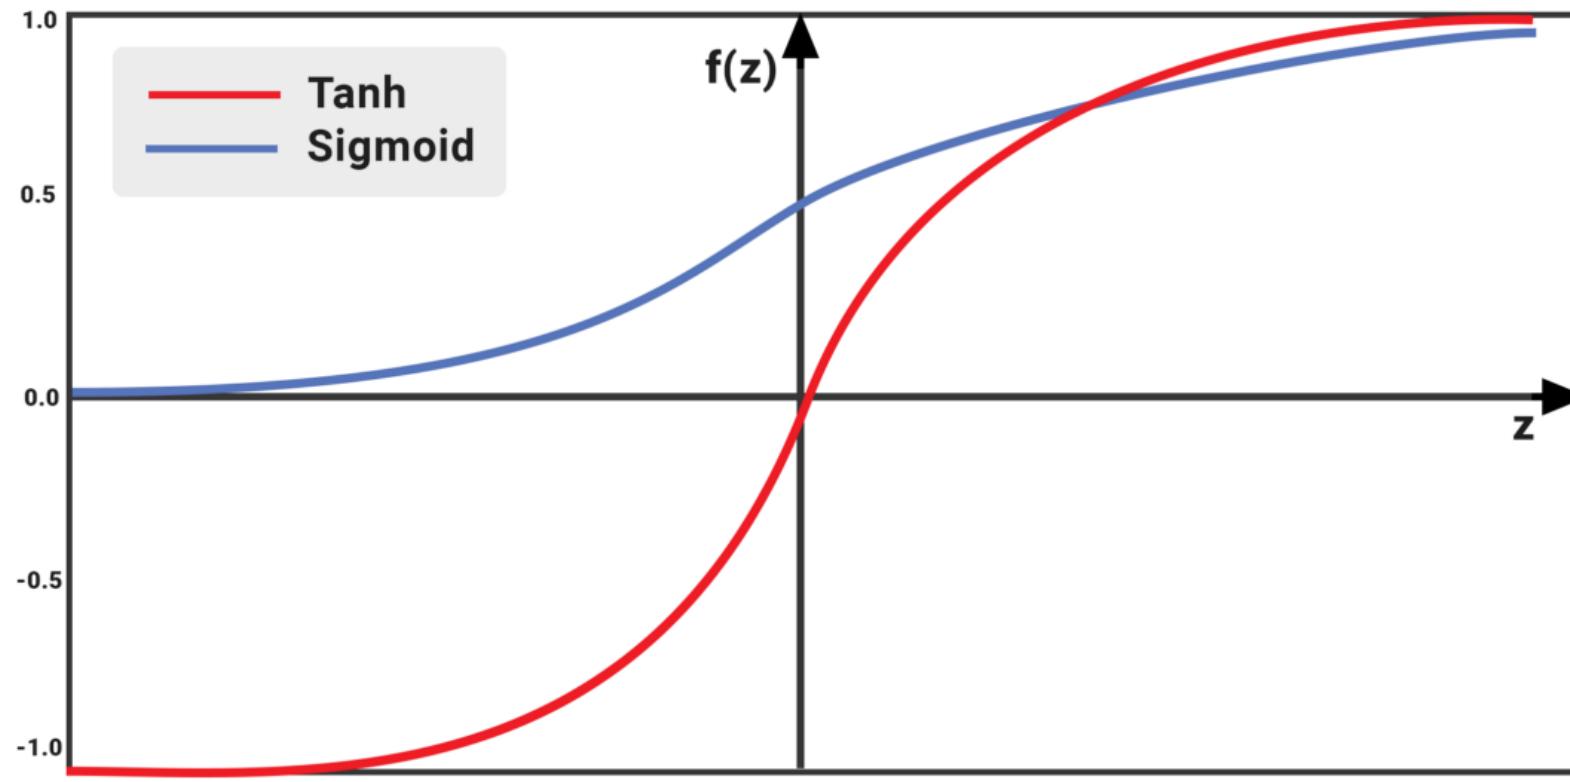
激活函数与GPU

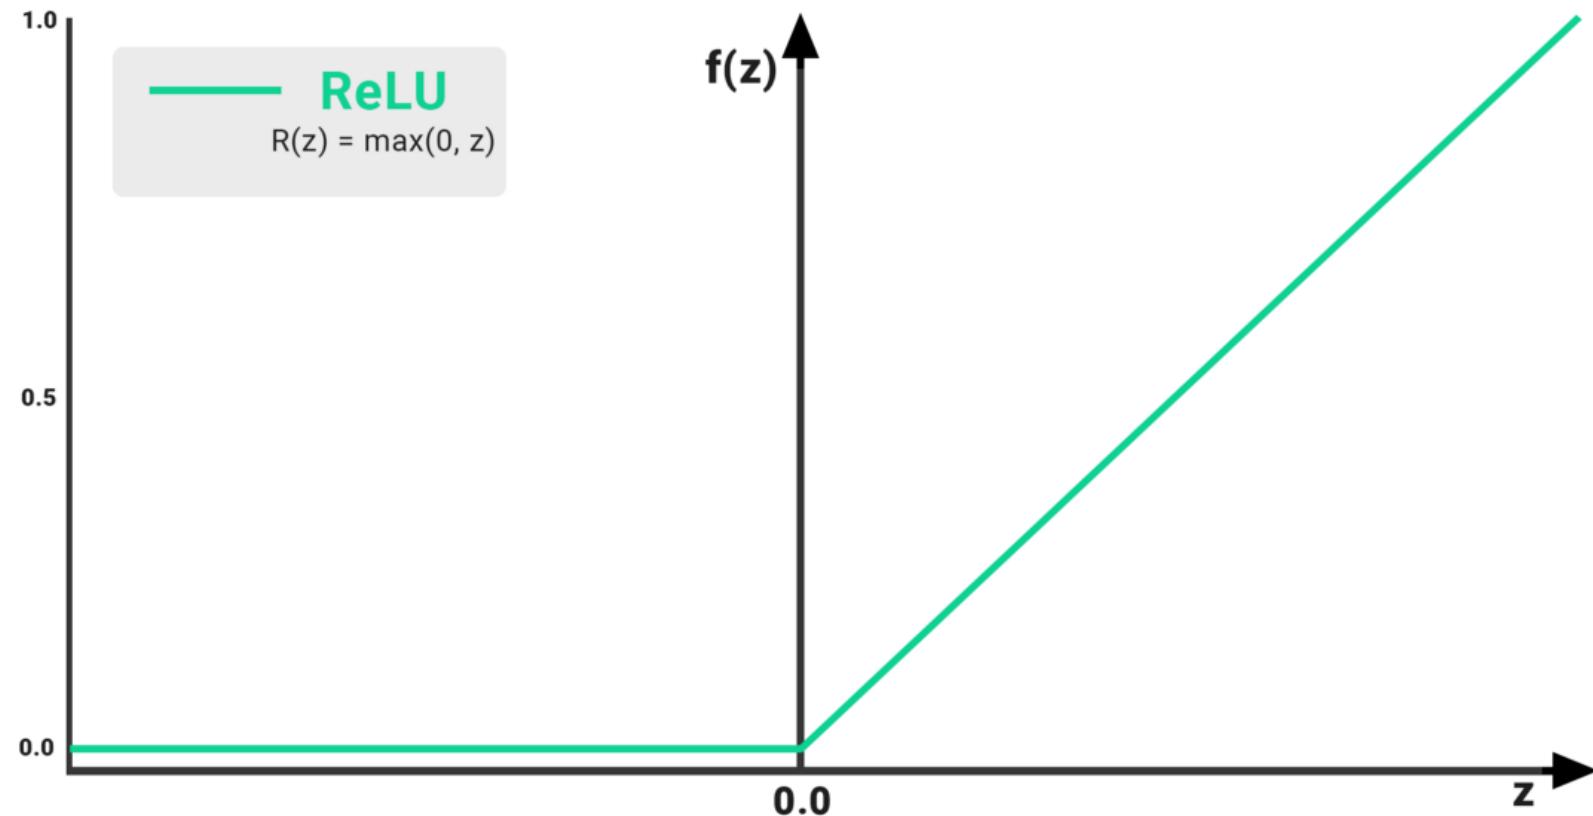
Thank You.



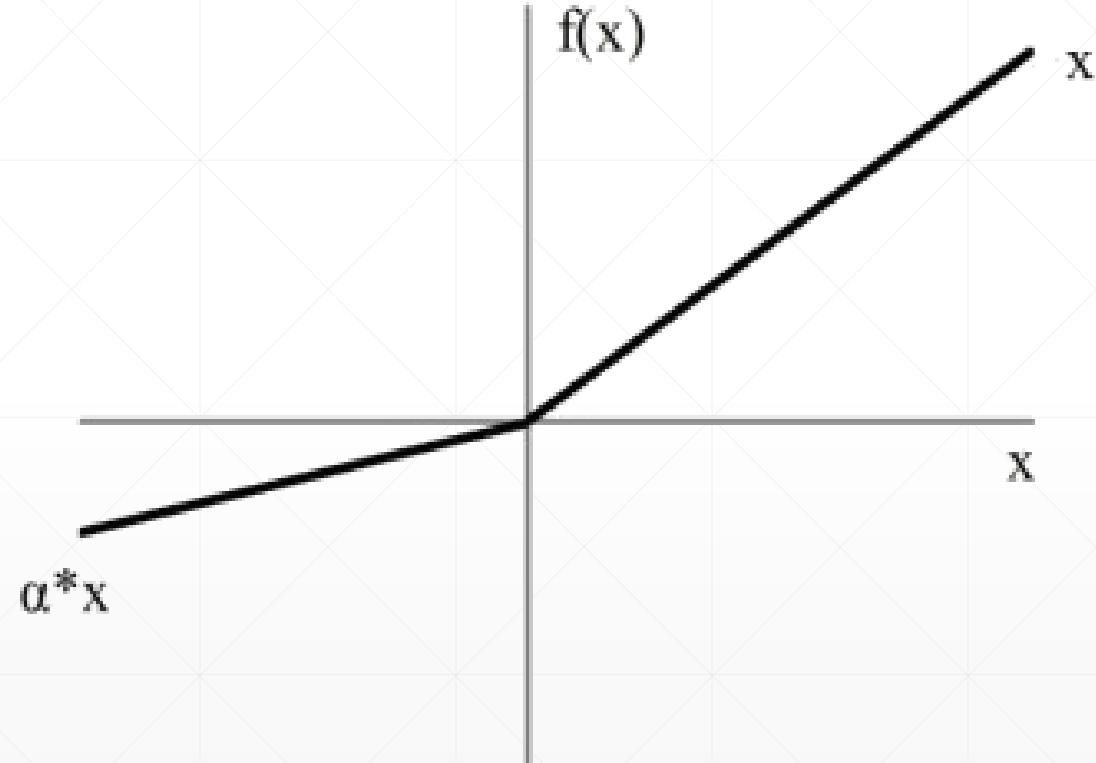
激活函数与GPU加速

主讲人：龙良曲





Leaky ReLU



simply

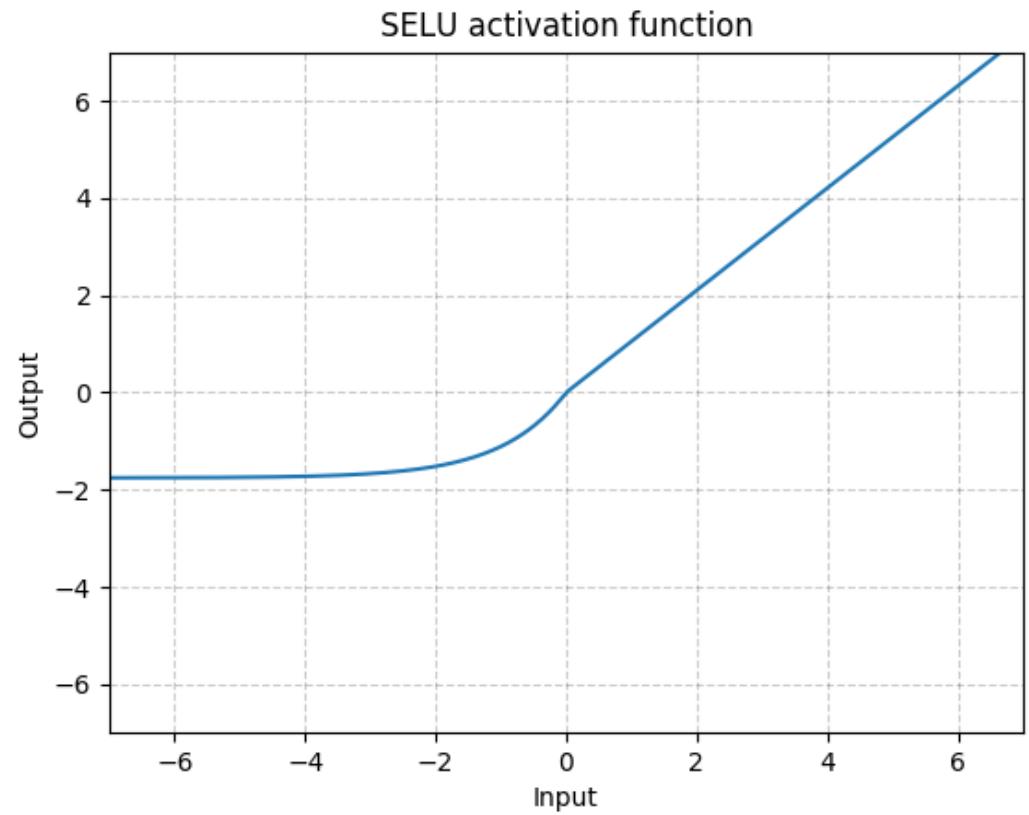


```
self.model = nn.Sequential(  
    nn.Linear(784, 200),  
    nn.LeakyReLU(inplace=True),  
    nn.Linear(200, 200),  
    nn.LeakyReLU(inplace=True),  
    nn.Linear(200, 10),  
    nn.LeakyReLU(inplace=True),  
)
```

SELU

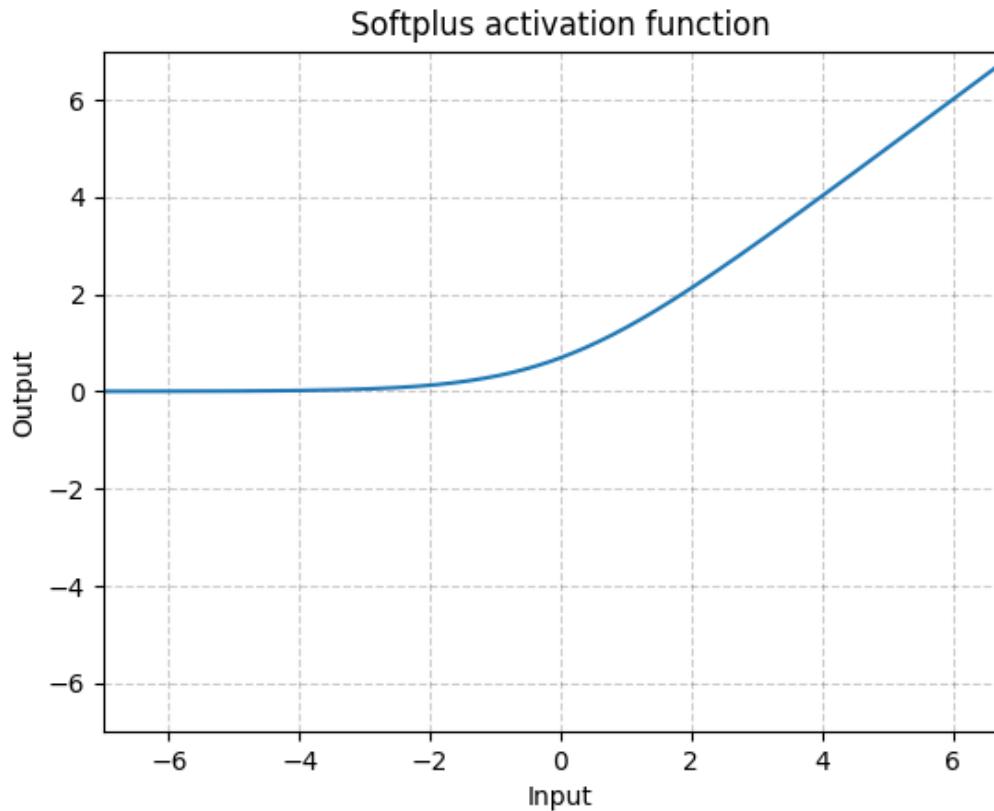
$$\text{SELU}(x) = \text{scale} * (\max(0, x) + \min(0, \alpha * (\exp(x) - 1)))$$

with $\alpha = 1.6732632423543772848170429916717$ and
scale = 1.0507009873554804934193349852946.



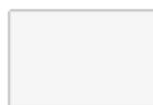
softplus

$$\text{Softplus}(x) = \frac{1}{\beta} * \log(1 + \exp(\beta * x))$$



GPU accelerated

```
● ● ●  
  
device = torch.device('cuda:0')  
net = MLP().to(device)  
optimizer = optim.SGD(net.parameters(), lr=learning_rate)  
criterion = nn.CrossEntropyLoss().to(device)  
  
for epoch in range(epochs):  
  
    for batch_idx, (data, target) in enumerate(train_loader):  
        data = data.view(-1, 28*28)  
        data, target = data.to(device), target.cuda()
```



Ethernet
Not connected



Ethernet
S: 0 R: 0 Kbps



Ethernet
S: 0 R: 0 Kbps



Wi-Fi
S: 0 R: 0 Kbps



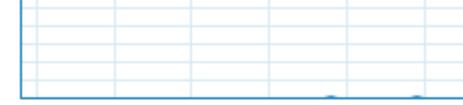
Bluetooth PAN
Not connected



GPU 0
Intel(R) HD Graphics
0%

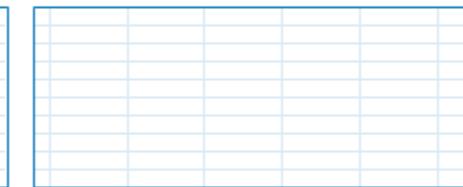
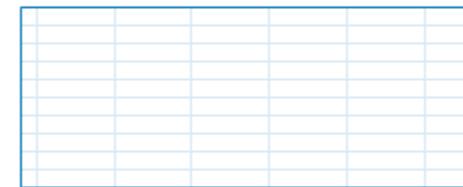


GPU 1
NVIDIA GeForce GTX
1%



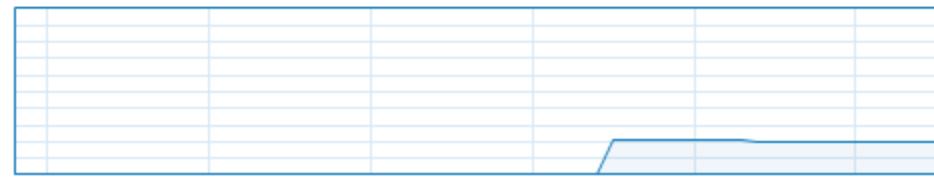
▼ Video Encode

0% ▼ Video Decode 0%



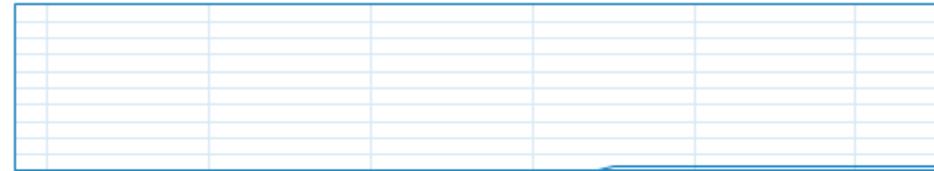
Dedicated GPU memory usage

4.0 GB



Shared GPU memory usage

3.9 GB



Utilization

1%

Dedicated GPU memory

0.8/4.0 GB

Driver version:

Driver date:

GPU Memory

0.9/7.9 GB

Shared GPU memory

0.1/3.9 GB

DirectX version:

Physical location:

Hardware reserved mem...
Physical location:
Hardware reserved mem...



Fewer details

Open Resource Monitor

下一课时

测试

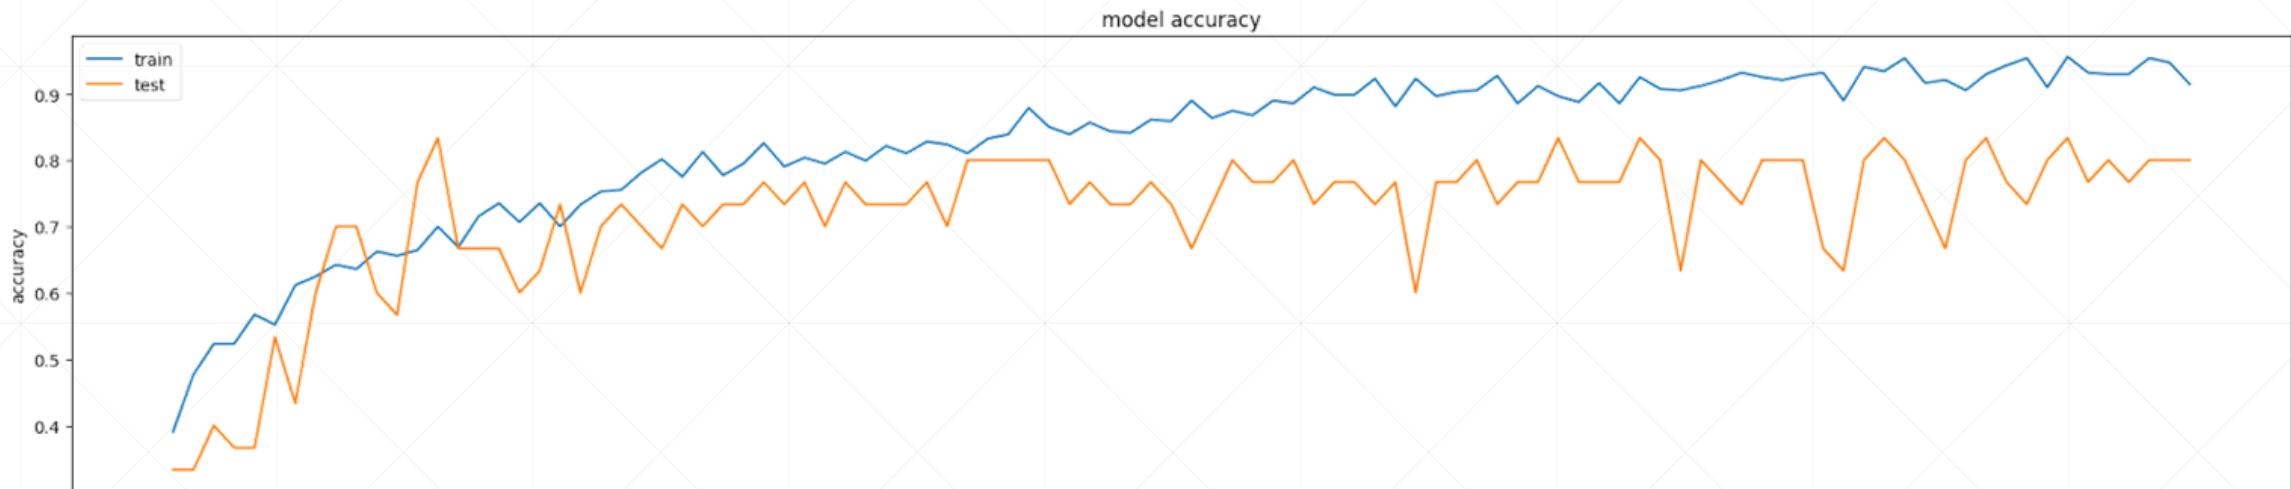
Thank You.



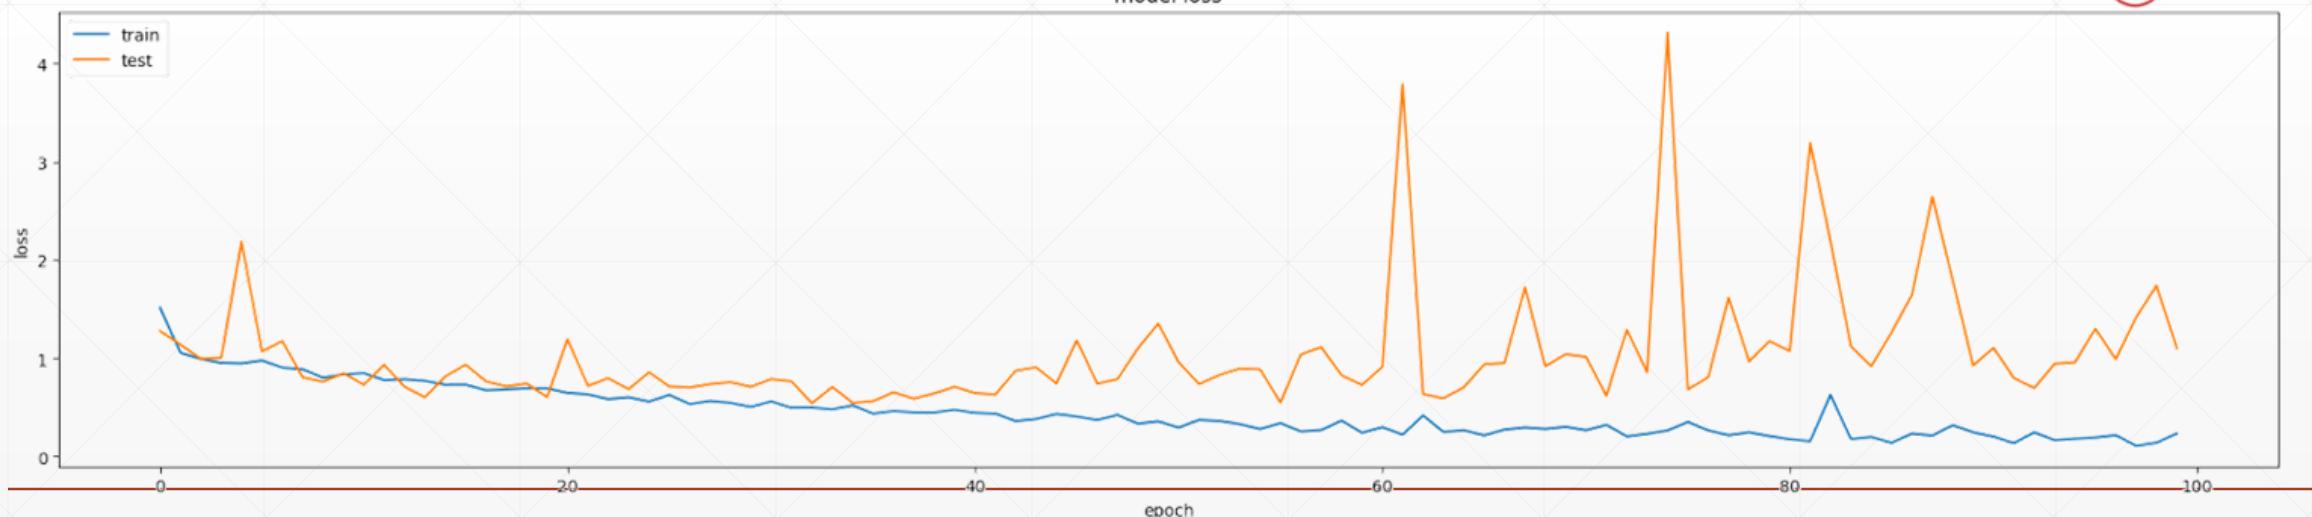
测试

主讲人：龙良曲

Loss != Accuracy



epoch
model loss



argmax



```
In [60]: logits=torch.rand(4, 10)
```

```
In [62]: pred=F.softmax(logits, dim=1)
```

```
In [63]: pred.shape
```

```
Out[63]: torch.Size([4, 10])
```

```
In [64]: pred_label=pred.argmax(dim=1)
```

```
In [65]: pred_label
```

```
Out[65]: tensor([9, 5, 9, 4])
```

```
In [66]: logits.argmax(dim=1)
```

```
Out[66]: tensor([9, 5, 9, 4])
```

```
In [67]: label=torch.tensor([9,3,2,4])
```

```
In [68]: correct=torch.eq(pred_label, label)
```

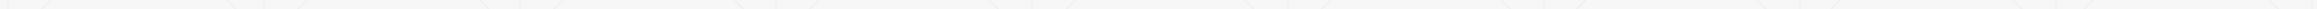
```
Out[69]: tensor([1, 0, 0, 1], dtype=torch.uint8)
```

```
In [71]: correct.sum().float().item()/4.
```

```
Out[71]: 0.5
```

When to test

- test once per several batch
- test once per epoch
- epoch V.S. step?





```
test_loss = 0
correct = 0
for data, target in test_loader:
    data = data.view(-1, 28 * 28)
    data, target = data.to(device), target.cuda()
    logits = net(data)
    test_loss += criterion(logits, target).item()

    pred = logits.argmax(dim=1)
    correct += pred.eq(target).float().sum().item()

test_loss /= len(test_loader.dataset)
print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)'.format(
    test_loss, correct, len(test_loader.dataset),
    100. * correct / len(test_loader.dataset)))
```

下一课时

Visdom可视化

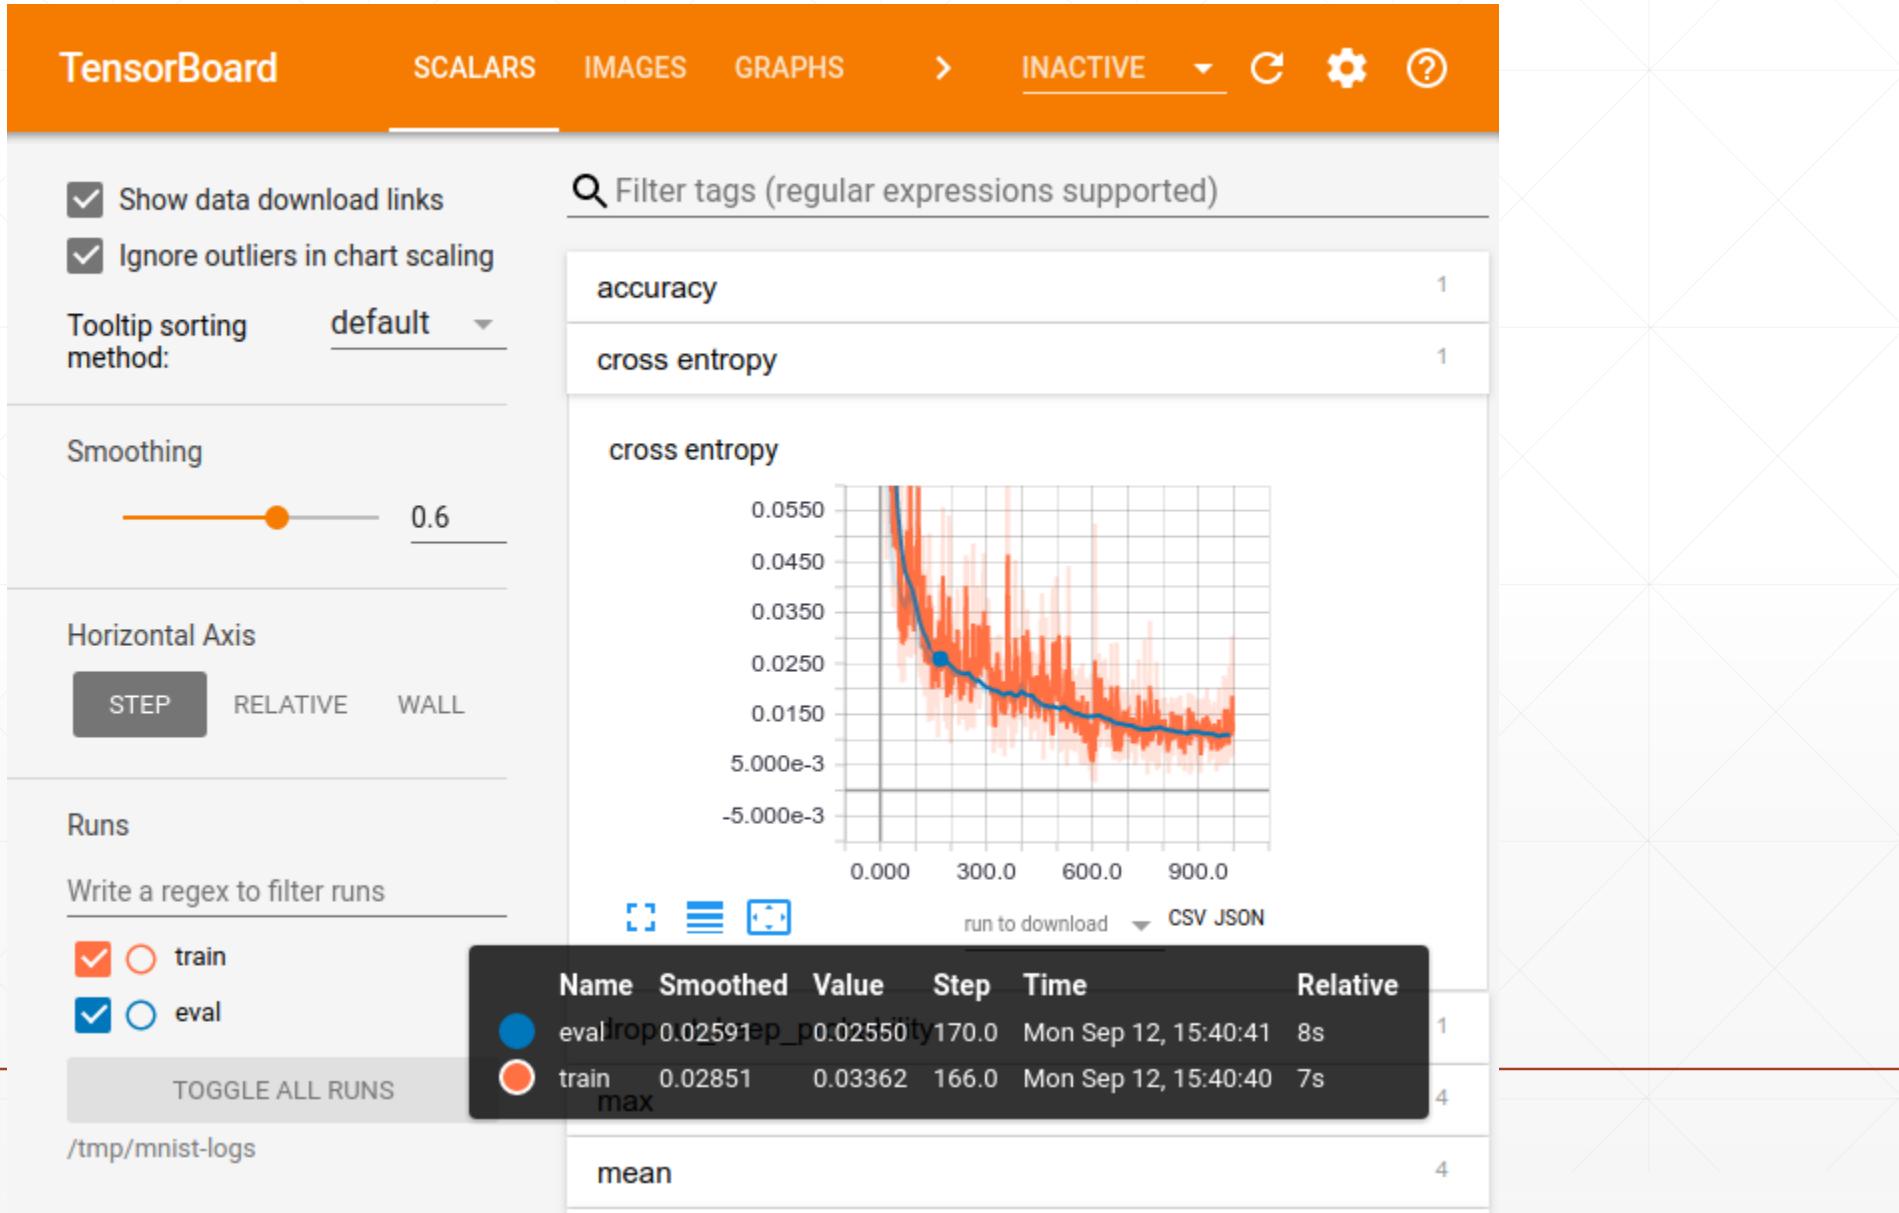
Thank You.



Visdom 可视化

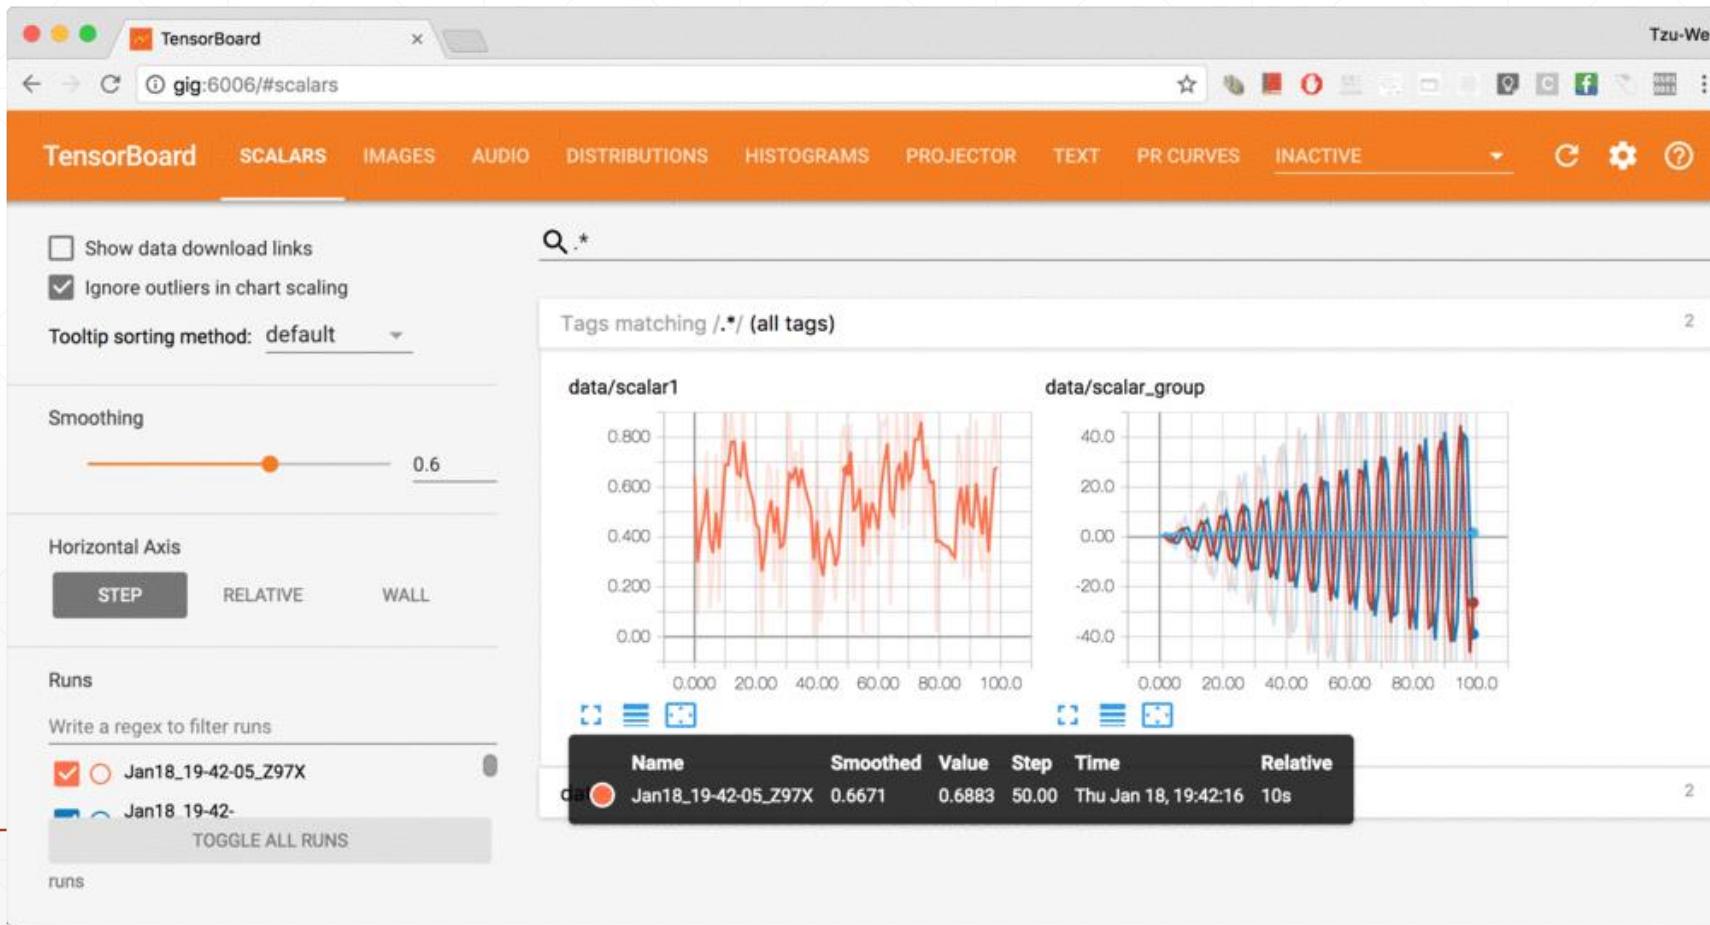
主讲人：龙良曲

TensorBoard?



TensorboardX

- pip install tensorboardX



TensorboardX



```
from tensorboardX import SummaryWriter

writer = SummaryWriter()
writer.add_scalar('data/scalar1', dummy_s1[0], n_iter)

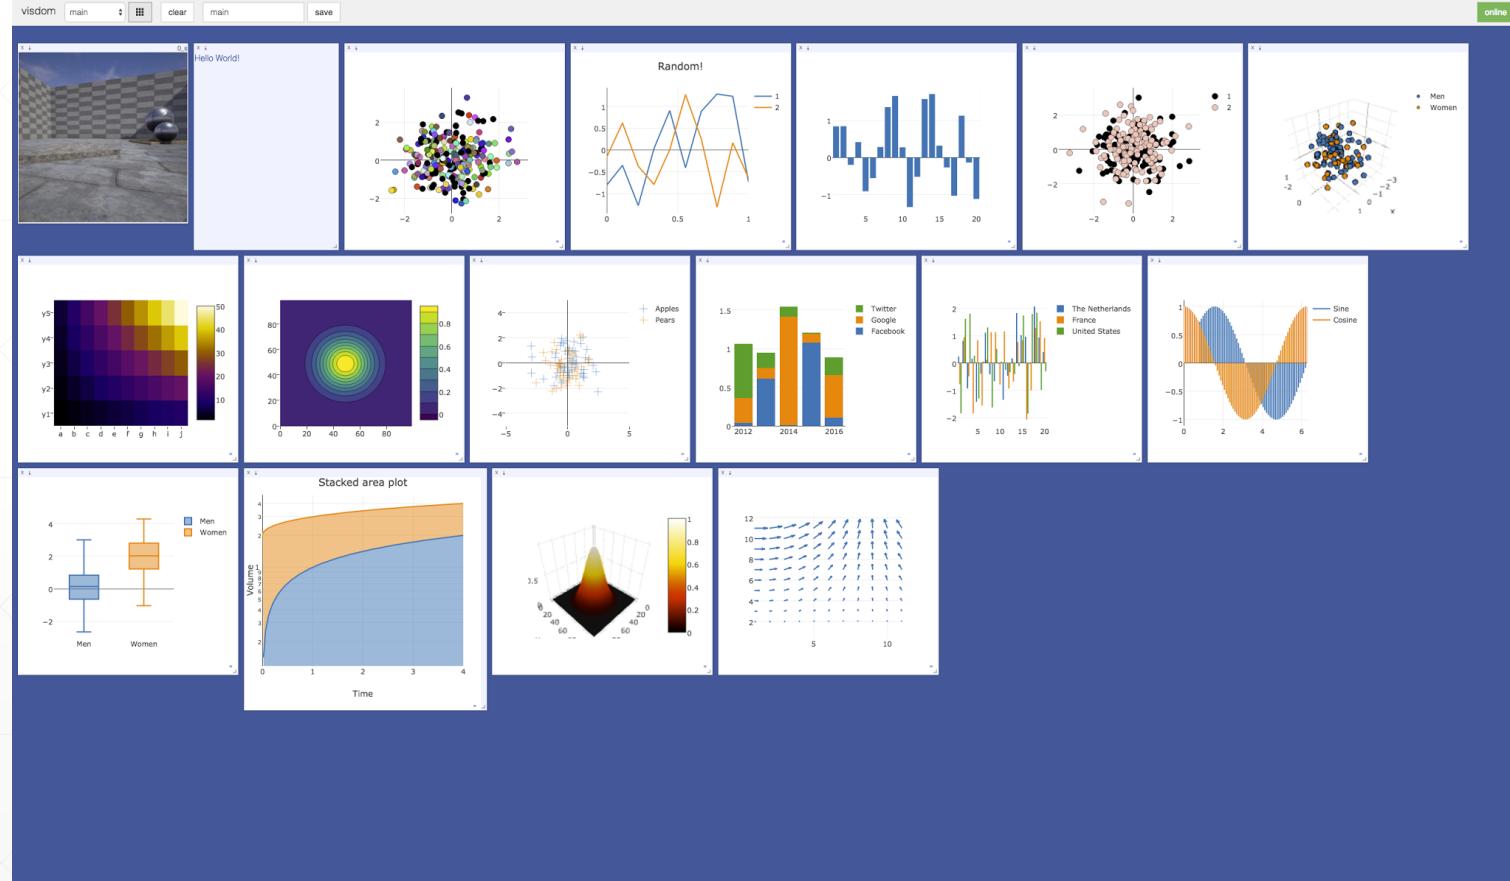
writer.add_scalars('data/scalar_group', {'xsinx': n_iter * np.sin(n_iter),
                                         'xcosx': n_iter * np.cos(n_iter),
                                         'arctanx': np.arctan(n_iter)}, n_iter)

writer.add_image('Image', x, n_iter)
writer.add_text('Text', 'text logged at step:' + str(n_iter), n_iter)

for name, param in resnet18.named_parameters():
    writer.add_histogram(name, param.clone().cpu().data.numpy(), n_iter)

writer.close()
```

Visdom from Facebook



Step 1. install

```
C:\Users\drage>pip install visdom
Collecting visdom
  Downloading https://files.pythonhosted.org/packages/c1/48/d90e1519768107811fd6e7760bea46fff9e9c9ffb490441684003ae634a9
  /visdom-0.1.8.5.tar.gz (248kB)
    100% |████████████████████████████████| 256kB 466kB/s
Requirement already satisfied: numpy>=1.8 in c:\programdata\conda\lib\site-packages (from visdom)
Requirement already satisfied: scipy in c:\programdata\conda\lib\site-packages (from visdom)
Requirement already satisfied: requests in c:\programdata\conda\lib\site-packages (from visdom)
Requirement already satisfied: tornado in c:\programdata\conda\lib\site-packages (from visdom)
Requirement already satisfied: pyzmq in c:\programdata\conda\lib\site-packages (from visdom)
Requirement already satisfied: six in c:\programdata\conda\lib\site-packages (from visdom)
Collecting torchfile (from visdom)
  Downloading https://files.pythonhosted.org/packages/91/af/5b305f86f2d218091af657ddb53f984ecbd9518ca9fe8ef4103a007252c9
  /torchfile-0.1.0.tar.gz
Collecting websocket-client (from visdom)
  Downloading https://files.pythonhosted.org/packages/26/2d/f749a5c82f6192d77ed061a38e02001afcba55fe8477336d26a950ab17ce
  /websocket_client-0.54.0-py2.py3-none-any.whl (200kB)
    100% |████████████████████████████████| 204kB 5.4MB/s
```

Step2. run server damon

```
C:\Users\drage>python -m visdom.server
Downloading scripts. It might take a while.
ERROR:root:Error 404 while downloading https://unpkg.com/layout-bin-packer@1.4.0
```

Step2. run server damon

```
C:\Users\drage>python -m visdom.server
Downloading scripts. It might take a while.
ERROR:root:Error 404 while downloading https://unpkg.com/layout-bin-packer@1.4.0
It's Alive!
INFO:root:Application Started
You can navigate to http://localhost:8097
INFO:tornado.access:200 POST /win_exists (::1) 0.00ms
INFO:tornado.access:200 POST /events (::1) 1.00ms
INFO:tornado.access:200 POST /win_exists (::1) 0.00ms
INFO:tornado.access:200 POST /update (::1) 1.00ms
INFO:tornado.access:200 POST /win_exists (::1) 0.00ms
INFO:tornado.access:200 POST /update (::1) 1.03ms
```

install from source

The screenshot shows a web browser window with multiple tabs open, including ones for Baidu, PyTorch, and PyTorch Tutorials. The main content is the GitHub repository for `facebookresearch/visdom`. The search bar at the top has the word "install" typed into it, and a dropdown menu shows "1/13" results. The repository page itself has a dark header with the GitHub logo and navigation links for Pull requests, Issues, Marketplace, and Explore. Below the header, the repository name `facebookresearch / visdom` is shown, along with metrics: 156 Watchers, 5,339 Stars, and 650 Forks. A description of the repository follows: "A flexible tool for creating, organizing, and sharing visualizations of live, rich data. Supports Torch and Numpy." The repository stats section shows 225 commits, 12 branches, 0 releases, and 63 contributors. A "View license" link is also present. The commit history lists several recent changes, such as "Changed code to parse hostname to use urllib (#531)" by `luisenp` and `facebook-github-bot`, and "Line different dash types (#523)" by `luisenp`. Other commits include "Make Image Pane scrolling browser independent (#526)", "Properties pane added", and "Prepares dependency versioning, completes 0.1.8". The bottom of the page includes standard GitHub repository controls like "Create new file", "Upload files", "Find file", and "Clone or download". The browser's taskbar at the bottom shows various pinned icons, and the system tray indicates the date and time as 12/31/2018 at 11:02 PM.

install
1/13

facebookresearch / visdom

Watch 156 Unstar 5,339 Fork 650

Code Issues 39 Pull requests 1 Projects 0 Insights

A flexible tool for creating, organizing, and sharing visualizations of live, rich data. Supports Torch and Numpy.

225 commits 12 branches 0 releases 63 contributors View license

Branch: master New pull request Create new file Upload files Find file Clone or download

luisenp and facebook-github-bot Changed code to parse hostname to use urllib (#531) ... Latest commit 9caf4dc 12 days ago

.github/ISSUE_TEMPLATE Create Issue Templates (#384) 6 months ago

example Line different dash types (#523) 17 days ago

js Make Image Pane scrolling browser independent (#526) 13 days ago

py/visdom Changed code to parse hostname to use urllib (#531) 12 days ago

th Properties pane added 7 months ago

.gitignore Prepares dependency versioning, completes 0.1.8 8 months ago

CODE_OF_CONDUCT.md Add code of conduct file (#488) 3 months ago

CONTRIBUTING.md Firewall update (#379) 7 months ago

LICENSE options -> opts, removing lua options parameter in favor of opts 8 months ago

11:02 PM 12/31/2018

lines: single trace



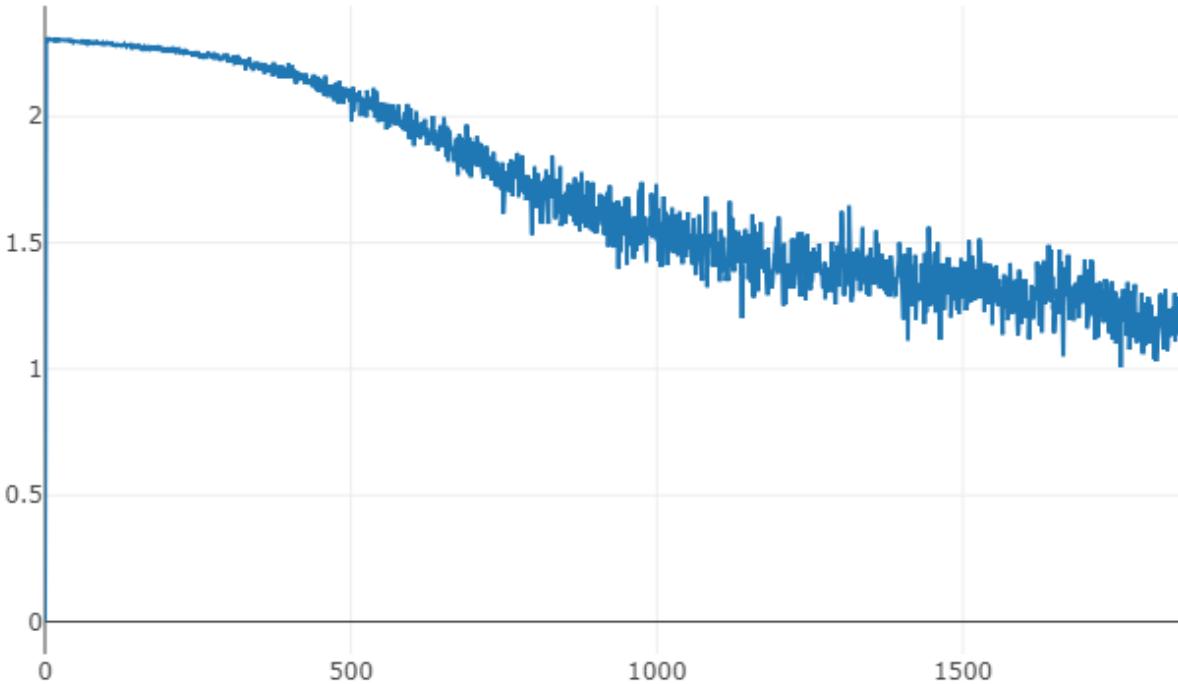
```
from visdom import Visdom

viz = Visdom()

viz.line([0.], [0.], win='train_loss', opts=dict(title='train loss'))

viz.line([loss.item()], [global_step], win='train_loss', update='append')
```

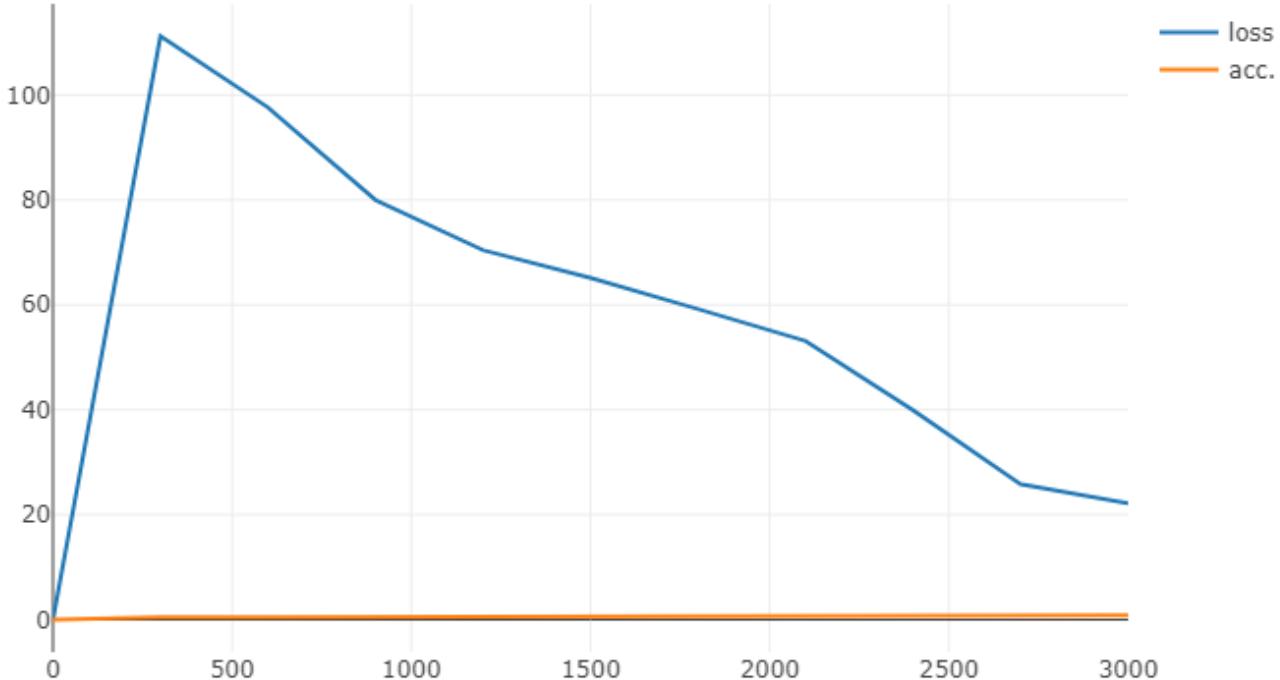
train loss



lines: multi-traces

```
● ● ●  
from visdom import Visdom  
  
viz = Visdom()  
viz.line([[0.0, 0.0]], [0.], win='test', opts=dict(title='test loss&acc.',  
                                              legend=['loss', 'acc.']))  
  
viz.line([[test_loss, correct / len(test_loader.dataset)]],  
        [global_step], win='test', update='append')
```

test loss&acc.



visual X

```
● ● ●  
from visdom import Visdom  
  
viz = Visdom()  
  
viz.images(data.view(-1, 1, 28, 28), win='x')  
viz.text(str(pred.detach().cpu().numpy()), win='pred',  
         opts=dict(title='pred'))
```

7	1	2	8	9	0	9	1
5	7	6	0	3	2	5	1
2	4	2	7	8	6	2	2
8	6	2	6	7	9	0	4
6	4	4	1	5	2	9	1
8	6	8	7	1	1	8	2
7	1	3	2	2	6	3	0
4	8	0	3	6	8	8	5
6	8	0	8	3	0	4	8
7	8	0	2	0	8	8	2
1	1	2	8	1	0	0	0

```
x & [70289091536032512427  
86228626790464428291  
869711827132263045076  
588680830489427701990  
753669314134864371527  
079519615441943067179  
467354100995911510876  
113940488472777288141  
93487003853700363392  
51045072086454]
```

下一课时

train-val-test

Thank You.
